

# Kutay\_Ugurlu\_EE583\_HW7

January 23, 2022

## 1 SETUP

```
[ ]: import torch
import numpy as np
import torchvision.datasets as datasets
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
from torchsummary import summary
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: #DEFINE YOUR DEVICE
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

print(device) #if cpu, go Runtime-> Change runtime type-> Hardware accelerator
↳GPU -> Save -> Redo previous steps
```

cuda:0

```
[ ]: #DOWNLOAD CIFAR-10 DATASET
train_data = datasets.CIFAR10('./data', train = True, download = True,
↳transform = transforms.ToTensor())

test_data = datasets.CIFAR10('./data', train = False, transform = transforms.
↳ToTensor())
```

```
[ ]: #DEFINE DATA GENERATOR
batch_size = 100
train_generator = torch.utils.data.DataLoader(train_data, batch_size =
↳batch_size, shuffle = True)
test_generator = torch.utils.data.DataLoader(test_data, batch_size =
↳batch_size, shuffle = False)
```

## 2 QUESTIONS

### 2.1 Q1

```
[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ1(torch.nn.Module):
    def __init__(self):
        super(CNNQ1, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 8, kernel_size = 4, stride = 1)
        self.conv2 = torch.nn.Conv2d(8, 16, kernel_size = 4, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(400, 256)
        self.fc2 = torch.nn.Linear(256, 64)
        self.fc3 = torch.nn.Linear(64, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
    def forward(self, x):
        hidden = self.mpool(self.relu(self.conv1(x)))
        hidden = self.mpool(self.relu(self.conv2(hidden)))
        hidden = hidden.view(-1,400)
        hidden = self.relu(self.fc1(hidden))
        hidden = self.relu(self.fc2(hidden))
        output = self.fc3(hidden)
        return output
```

```
[ ]: #CREATE MODEL
model = CNNQ1()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

# SUMMARY
summary(model, (3,32,32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 29, 29]	392
ReLU-2	[-1, 8, 29, 29]	0
MaxPool2d-3	[-1, 8, 14, 14]	0
Conv2d-4	[-1, 16, 11, 11]	2,064
ReLU-5	[-1, 16, 11, 11]	0
MaxPool2d-6	[-1, 16, 5, 5]	0

Linear-7	[-1, 256]	102,656
ReLU-8	[-1, 256]	0
Linear-9	[-1, 64]	16,448
ReLU-10	[-1, 64]	0
Linear-11	[-1, 10]	650

=====  
Total params: 122,210

Trainable params: 122,210

Non-trainable params: 0

-----  
Input size (MB): 0.01

Forward/backward pass size (MB): 0.15

Params size (MB): 0.47

Estimated Total Size (MB): 0.63  
-----

```
[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
        if (batch_idx+1) % batch_size == 0:
            print('Epoch: {}/{} [Batch: {}/{} ( {:.0f}%)]\tLoss: {:.6f}'.format(
                i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
                100. * (batch_idx+1) / len(train_generator), loss.item()))
```

Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 2.008565

Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 1.707132

Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 1.772573

Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 1.631163

Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 1.389537

Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 1.421544

Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 1.322189

Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 1.521783

Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 1.428849

Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 1.630560

```

Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 1.351266
Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 1.573170
Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 1.297113
Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 1.258028
Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 1.333989
Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 1.322998
Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 1.046606
Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 1.205253
Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 1.333849
Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 1.283796
Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 1.230737
Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 1.080133
Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 1.232069
Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 1.105746
Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 1.262319
Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 1.283407
Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 1.138714
Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 1.239647
Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 0.980552
Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 1.403119
Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 1.114779
Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 1.089277
Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 1.201144
Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 1.393427
Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 1.164613
Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 1.267781
Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 1.099852
Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 1.168808
Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 1.133110
Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 0.952967
Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 1.156250
Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 0.927581
Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 1.209802
Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 1.319026
Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 0.907355
Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 1.167614
Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 1.059658
Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 0.972945
Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 1.021400
Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 0.903114


```

```

[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')

```

```
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q1.png')
```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1

print('Validation accuracy: %.2f%%' %((100*correct)//(total)))
```

Validation accuracy: 62.00%

## 2.2 Q2

### 2.2.1 Optimizer changed to SGD.

```
[ ]: #CREATE MODEL
model = CNNQ1()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.01

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr = learning_rate, momentum=0.
↪9)

# SUMMARY
summary(model, (3,32,32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 29, 29]	392
ReLU-2	[-1, 8, 29, 29]	0
MaxPool2d-3	[-1, 8, 14, 14]	0
Conv2d-4	[-1, 16, 11, 11]	2,064
ReLU-5	[-1, 16, 11, 11]	0
MaxPool2d-6	[-1, 16, 5, 5]	0
Linear-7	[-1, 256]	102,656
ReLU-8	[-1, 256]	0
Linear-9	[-1, 64]	16,448
ReLU-10	[-1, 64]	0
Linear-11	[-1, 10]	650

=====  
Total params: 122,210  
Trainable params: 122,210  
Non-trainable params: 0  
=====

Input size (MB): 0.01  
Forward/backward pass size (MB): 0.15  
Params size (MB): 0.47  
Estimated Total Size (MB): 0.63  
=====

```
[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)
```

```

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):

    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
    if (batch_idx+1) % batch_size == 0:
        print('Epoch: {}/{} [Batch: {}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
            i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
            100. * (batch_idx+1) / len(train_generator), loss.item()))

```

```

Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 0.736044
Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 0.781371
Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 0.696068
Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 0.812954
Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 0.862274
Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 0.966946
Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 0.963485
Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 0.817716
Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 0.972773
Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 0.975334
Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 0.818820
Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 0.746322
Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 0.679832
Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 0.798428
Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 0.838489
Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 0.807739
Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 0.734484
Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 0.553443
Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 0.861763
Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 1.028033
Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 0.700818
Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 0.788175
Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 0.722639
Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 0.818164
Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 0.806038
Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 0.674392
Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 0.644142
Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 0.570559
Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 0.622261

```

```

Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 0.609262
Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 0.668045
Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 0.665257
Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 0.586246
Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 0.847288
Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 0.602923
Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 0.470082
Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 0.598713
Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 0.573132
Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 0.541481
Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 0.550917
Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 0.549263
Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 0.645425
Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 0.650120
Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 0.588912
Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 0.655012
Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 0.501349
Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 0.421503
Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 0.499366
Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 0.637380
Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 0.617409


```

```

[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q2.png')

```





```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1

print('Validation accuracy: %.2f%%' %((100*correct)//(total)))
```

Validation accuracy: 61.00%

With LR = 0.01 and Momentum = 0.9, the validation accuracy was lower than that of Adam's. Hence, proceed with Adam.

## 2.3 Q3

### 2.3.1 Activation changed to sigmoid.

```
[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ3(torch.nn.Module):
    def __init__(self):
        super(CNNQ3, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 8, kernel_size = 4, stride = 1)
        self.conv2 = torch.nn.Conv2d(8, 16, kernel_size = 4, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(400, 256)
        self.fc2 = torch.nn.Linear(256, 64)
        self.fc3 = torch.nn.Linear(64, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
    def forward(self, x):
        hidden = self.mpool(self.sigmoid(self.conv1(x)))
        hidden = self.mpool(self.sigmoid(self.conv2(hidden)))
        hidden = hidden.view(-1,400)
        hidden = self.sigmoid(self.fc1(hidden))
        hidden = self.sigmoid(self.fc2(hidden))
        output = self.fc3(hidden)
        return output
```

```
[ ]: #CREATE MODEL
model = CNNQ3()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

# SUMMARY
summary(model, (3,32,32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 29, 29]	392
Sigmoid-2	[-1, 8, 29, 29]	0
MaxPool2d-3	[-1, 8, 14, 14]	0
Conv2d-4	[-1, 16, 11, 11]	2,064
Sigmoid-5	[-1, 16, 11, 11]	0
MaxPool2d-6	[-1, 16, 5, 5]	0

Linear-7	[-1, 256]	102,656
Sigmoid-8	[-1, 256]	0
Linear-9	[-1, 64]	16,448
Sigmoid-10	[-1, 64]	0
Linear-11	[-1, 10]	650

=====  
Total params: 122,210

Trainable params: 122,210

Non-trainable params: 0  
-----

Input size (MB): 0.01

Forward/backward pass size (MB): 0.15

Params size (MB): 0.47

Estimated Total Size (MB): 0.63  
-----

```
[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
    if (batch_idx+1) % batch_size == 0:
        print('Epoch: {}/{} [Batch: {}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
            i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
            100. * (batch_idx+1) / len(train_generator), loss.item()))
```

Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 2.304116

Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 2.163043

Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 2.151380

Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 2.033580

Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 1.966078

Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 1.993485

Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 2.116996

Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 2.015048

Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 2.021232

Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 1.953170

```

Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 1.814739
Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 1.778557
Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 1.842888
Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 1.768866
Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 1.881565
Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 1.762808
Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 1.810751
Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 1.908579
Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 1.915297
Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 1.596877
Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 1.747207
Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 1.686767
Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 1.713181
Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 1.652671
Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 1.728626
Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 1.612773
Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 1.728046
Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 1.881076
Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 1.758918
Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 1.678548
Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 1.635035
Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 1.580833
Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 1.826030
Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 1.701180
Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 1.635299
Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 1.600815
Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 1.398542
Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 1.677052
Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 1.696446
Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 1.707719
Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 1.536893
Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 1.645865
Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 1.622025
Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 1.665705
Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 1.710800
Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 1.508257
Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 1.513326
Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 1.608315
Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 1.468388
Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 1.439004


```

```

[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')

```

```
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q3.png')
```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1

print('Validation accuracy: %.2f%%' %((100*correct)//(total)))
```

Validation accuracy: 45.00%

The network with activations Sigmoid, rather than ReLU, performed remarkably poorer, both in terms of training loss and validation accuracy. Returning back to ReLU.

## 2.4 Q4

### 2.4.1 Increase the kernel size of convolutional layers

```
[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ4(torch.nn.Module):
    def __init__(self):
        super(CNNQ4, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 8, kernel_size = 6, stride = 1)
        self.conv2 = torch.nn.Conv2d(8, 16, kernel_size = 6, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(256, 256)
        self.fc2 = torch.nn.Linear(256, 64)
        self.fc3 = torch.nn.Linear(64, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
    def forward(self, x):
        hidden = self.mpool(self.relu(self.conv1(x)))
        hidden = self.mpool(self.relu(self.conv2(hidden)))
        hidden = hidden.view(-1,256)
        hidden = self.relu(self.fc1(hidden))
        hidden = self.relu(self.fc2(hidden))
        output = self.fc3(hidden)
        return output
```

```
[ ]: #CREATE MODEL
model = CNNQ4()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

# SUMMARY
summary(model, (3,32,32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 27, 27]	872
ReLU-2	[-1, 8, 27, 27]	0

MaxPool2d-3	[-1, 8, 13, 13]	0
Conv2d-4	[-1, 16, 8, 8]	4,624
ReLU-5	[-1, 16, 8, 8]	0
MaxPool2d-6	[-1, 16, 4, 4]	0
Linear-7	[-1, 256]	65,792
ReLU-8	[-1, 256]	0
Linear-9	[-1, 64]	16,448
ReLU-10	[-1, 64]	0
Linear-11	[-1, 10]	650

=====

Total params: 88,386  
Trainable params: 88,386  
Non-trainable params: 0

-----

Input size (MB): 0.01  
Forward/backward pass size (MB): 0.12  
Params size (MB): 0.34  
Estimated Total Size (MB): 0.47

-----

```
[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
        if (batch_idx+1) % batch_size == 0:
            print('Epoch: {}/{} [Batch: {}/{} ( {:.0f}%)]\tLoss: {:.6f}'.format(
                i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
                100. * (batch_idx+1) / len(train_generator), loss.item()))
```

```
Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 2.021697
Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 1.766634
Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 1.758830
Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 1.860762
Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 1.632864
Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 1.545504
```

```

Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 1.519792
Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 1.579289
Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 1.519796
Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 1.484550
Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 1.619676
Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 1.521780
Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 1.455799
Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 1.505249
Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 1.304922
Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 1.463613
Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 1.270110
Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 1.268965
Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 1.494770
Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 1.372415
Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 1.292886
Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 1.322384
Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 1.188125
Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 1.202625
Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 1.351026
Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 1.120500
Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 1.478141
Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 1.217769
Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 1.172160
Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 1.160467
Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 1.335635
Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 1.267392
Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 1.193899
Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 1.209125
Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 1.096546
Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 1.086738
Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 1.088318
Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 1.057433
Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 0.982501
Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 1.267817
Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 1.046066
Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 1.087351
Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 1.002679
Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 1.069166
Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 1.139246
Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 1.017682
Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 1.018468
Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 0.932353
Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 1.102373
Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 1.005971

```


```

[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)

```



```
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q4.png')
```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1
```

```
print('Validation accuracy: %.2f%%' %((100*correct)/(total)))
```

Validation accuracy: 57.00%

Increasing the kernel size decreased accuracy with respect to the Q1 network with Adams and Kernel size 4. This may be attributed to the decreased number of parameters in the network, due to the increased kernel size resulting in shrinked outputs, hence less number of parameters in the fully connected layers. Returning back to kernel size 4.

## 2.5 Q5

### 2.5.1 Remove max pooling layers

```
[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ5(torch.nn.Module):
    def __init__(self):
        super(CNNQ5, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 8, kernel_size = 4, stride = 1)
        self.conv2 = torch.nn.Conv2d(8, 16, kernel_size = 4, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(10816, 256)
        self.fc2 = torch.nn.Linear(256, 64)
        self.fc3 = torch.nn.Linear(64, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
    def forward(self, x):
        hidden = (self.relu(self.conv1(x)))
        hidden = (self.relu(self.conv2(hidden)))
        hidden = hidden.view(-1,10816)
        hidden = self.relu(self.fc1(hidden))
        hidden = self.relu(self.fc2(hidden))
        output = self.fc3(hidden)
        return output
```

```
[ ]: #CREATE MODEL
model = CNNQ5()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

# SUMMARY
```

```
summary(model, (3,32,32))
```

```
-----
              Layer (type)              Output Shape          Param #
=====
              Conv2d-1              [-1, 8, 29, 29]           392
              ReLU-2              [-1, 8, 29, 29]            0
              Conv2d-3              [-1, 16, 26, 26]         2,064
              ReLU-4              [-1, 16, 26, 26]            0
              Linear-5              [-1, 256]         2,769,152
              ReLU-6              [-1, 256]            0
              Linear-7              [-1, 64]          16,448
              ReLU-8              [-1, 64]            0
              Linear-9              [-1, 10]            650
=====
Total params: 2,788,706
Trainable params: 2,788,706
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 0.27
Params size (MB): 10.64
Estimated Total Size (MB): 10.92
-----
```

```
[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)


loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
        if (batch_idx+1) % batch_size == 0:
            print('Epoch: {}/{} [Batch: {}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
                100. * (batch_idx+1) / len(train_generator), loss.item()))
```

Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 0.397587  
 Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 0.354198  
 Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 0.419669  
 Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 0.300957  
 Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 0.445359  
 Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 0.202522  
 Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 0.258415  
 Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 0.422641  
 Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 0.271738  
 Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 0.308269  
 Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 0.104597  
 Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 0.201747  
 Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 0.166502  
 Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 0.178280  
 Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 0.145707  
 Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 0.165880  
 Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 0.080247  
 Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 0.093291  
 Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 0.222174  
 Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 0.188313  
 Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 0.230045  
 Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 0.047383  
 Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 0.090266  
 Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 0.143519  
 Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 0.215926  
 Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 0.120997  
 Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 0.066584  
 Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 0.156069  
 Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 0.137690  
 Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 0.100795  
 Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 0.096036  
 Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 0.028743  
 Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 0.176817  
 Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 0.065263  
 Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 0.080358  
 Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 0.039991  
 Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 0.123598  
 Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 0.029759  
 Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 0.063611  
 Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 0.102855  
 Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 0.053693  
 Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 0.025046  
 Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 0.035936  
 Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 0.293280  
 Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 0.163612  
 Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 0.040707  
 Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 0.078210  
 Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 0.033172

Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 0.108378

Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 0.086482

```
[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q5.png')
```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)
```

```

for i in range(y_pred.shape[0]):
    if y_val[i]==y_pred[i]:
        correct += 1
    total +=1

print('Validation accuracy: %.2f%%' %((100*correct)/(total)))

```

Validation accuracy: 59.00%

Removing the max pool layers resulted in a significant decrease in training loss due to the overfitting caused from increased number of network parameters with a slight decrease in the validation accuracy. This is somehow parallel with the expectations since one would expect the network would perform poorer in the absence of decreased nonlinearity of max pool layers.

## 2.6 Q6

### 2.6.1 A convolutional layer is added.

```

[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ6(torch.nn.Module):
    def __init__(self):
        super(CNNQ6, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 8, kernel_size = 4, stride = 1)
        self.conv2 = torch.nn.Conv2d(8, 16, kernel_size = 4, stride = 1)
        self.conv3 = torch.nn.Conv2d(16, 32, kernel_size = 4, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(32, 256)
        self.fc2 = torch.nn.Linear(256, 64)
        self.fc3 = torch.nn.Linear(64, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
    def forward(self, x):
        hidden = self.mpool(self.relu(self.conv1(x)))
        hidden = self.mpool(self.relu(self.conv2(hidden)))
        hidden = self.mpool(self.relu(self.conv3(hidden)))
        hidden = hidden.view(-1,32)
        hidden = self.relu(self.fc1(hidden))
        hidden = self.relu(self.fc2(hidden))
        output = self.fc3(hidden)
        return output

```

```

[ ]: #CREATE MODEL
model = CNNQ6()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER

```

```

learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

# SUMMARY
summary(model, (3,32,32))

```

```

-----
Layer (type)           Output Shape          Param #
=====
      Conv2d-1         [-1, 8, 29, 29]        392
      ReLU-2           [-1, 8, 29, 29]         0
    MaxPool2d-3        [-1, 8, 14, 14]         0
      Conv2d-4         [-1, 16, 11, 11]       2,064
      ReLU-5           [-1, 16, 11, 11]         0
    MaxPool2d-6        [-1, 16, 5, 5]          0
      Conv2d-7         [-1, 32, 2, 2]         8,224
      ReLU-8           [-1, 32, 2, 2]          0
    MaxPool2d-9        [-1, 32, 1, 1]          0
      Linear-10         [-1, 256]              8,448
      ReLU-11           [-1, 256]               0
      Linear-12         [-1, 64]              16,448
      ReLU-13           [-1, 64]               0
      Linear-14         [-1, 10]               650
=====
Total params: 36,226
Trainable params: 36,226
Non-trainable params: 0
-----

Input size (MB): 0.01
Forward/backward pass size (MB): 0.15
Params size (MB): 0.14
Estimated Total Size (MB): 0.30
-----

```

```

[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()

```

```

y_pred = model(x_train)
loss = loss_fun(y_pred, y_train)
loss_values[num_of_batch*i+batch_idx] = loss.item()
loss.backward()
optimizer.step()
if (batch_idx+1) % batch_size == 0:
    print('Epoch: {}/{ } [Batch: {}/{ } ({:.0f}%)]\tLoss: {:.6f}'.format(
        i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
↪dataset),
        100. * (batch_idx+1) / len(train_generator), loss.item()))

```

```

Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 2.226628
Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 1.937245
Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 1.870508
Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 1.777901
Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 1.813008
Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 1.651842
Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 1.741103
Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 1.607360
Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 1.631606
Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 1.629384
Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 1.708601
Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 1.505062
Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 1.473297
Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 1.248842
Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 1.375444
Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 1.539607
Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 1.452272
Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 1.448867
Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 1.326008
Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 1.528908
Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 1.432323
Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 1.216753
Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 1.533282
Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 1.422818
Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 1.409254
Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 1.445814
Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 1.342880
Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 1.525958
Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 1.431715
Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 1.316404
Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 1.144156
Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 1.118805
Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 1.267522
Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 1.151852
Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 1.290665
Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 1.392779

```



```


Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 1.144426
Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 1.142717
Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 1.247062
Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 1.453298
Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 1.123607
Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 1.485599
Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 1.229701
Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 1.365828
Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 1.246860
Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 1.442385
Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 1.284320
Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 1.314914
Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 1.070150
Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 1.366385

```

```

[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q6.png')

```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1

print('Validation accuracy: %.2f%%' %((100*correct)//(total)))
```

Validation accuracy: 55.00%

Previous experiment showed that pooling layers are useful for the network. Adding one extra convolutional layer increased the number of layers in the network, hence would normally increase the capability of fitting more complex functions. However, with this 3<sup>rd</sup> convolutional layer, the weights of the network significantly decreased to half, hence we cannot deduce a meaningful conclusion from this experiment. Instead one can try more convolutional layers and more units in FCLs.

## 2.7 Q7

### 2.7.1 Increase the output size of the first convolutional layer

```
[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ7(torch.nn.Module):
    def __init__(self):
        super(CNNQ7, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 32, kernel_size = 4, stride = 1)
        self.conv2 = torch.nn.Conv2d(32, 64, kernel_size = 4, stride = 1)
        self.conv3 = torch.nn.Conv2d(64, 32, kernel_size = 4, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(128, 1024)
        self.fc2 = torch.nn.Linear(1024, 256)
        self.fc3 = torch.nn.Linear(256, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
    def forward(self, x):
        hidden = self.mpool(self.relu(self.conv1(x)))
```

```

hidden = self.mpool(self.relu(self.conv2(hidden)))
hidden = (self.relu(self.conv3(hidden)))
hidden = hidden.view(-1,128)
hidden = self.relu(self.fc1(hidden))
hidden = self.relu(self.fc2(hidden))
output = self.fc3(hidden)
return output

```

```

[ ]: #CREATE MODEL
model = CNNQ7()
model.to(device)

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

# SUMMARY
summary(model, (3,32,32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 29, 29]	1,568
ReLU-2	[-1, 32, 29, 29]	0
MaxPool2d-3	[-1, 32, 14, 14]	0
Conv2d-4	[-1, 64, 11, 11]	32,832
ReLU-5	[-1, 64, 11, 11]	0
MaxPool2d-6	[-1, 64, 5, 5]	0
Conv2d-7	[-1, 32, 2, 2]	32,800
ReLU-8	[-1, 32, 2, 2]	0
Linear-9	[-1, 1024]	132,096
ReLU-10	[-1, 1024]	0
Linear-11	[-1, 256]	262,400
ReLU-12	[-1, 256]	0
Linear-13	[-1, 10]	2,570

Total params: 464,266  
 Trainable params: 464,266  
 Non-trainable params: 0

Input size (MB): 0.01  
 Forward/backward pass size (MB): 0.61  
 Params size (MB): 1.77  
 Estimated Total Size (MB): 2.39

```
[ ]: #TRAIN THE MODEL
model.train()
epoch = 10

num_of_batch=np.int(len(train_generator.dataset)/batch_size)

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
        if (batch_idx+1) % batch_size == 0:
            print('Epoch: {}/{} [Batch: {}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
                100. * (batch_idx+1) / len(train_generator), loss.item()))
```

```
Epoch: 1/10 [Batch: 10000/50000 (20%)] Loss: 1.867347
Epoch: 1/10 [Batch: 20000/50000 (40%)] Loss: 1.461807
Epoch: 1/10 [Batch: 30000/50000 (60%)] Loss: 1.482144
Epoch: 1/10 [Batch: 40000/50000 (80%)] Loss: 1.393044
Epoch: 1/10 [Batch: 50000/50000 (100%)] Loss: 1.518363
Epoch: 2/10 [Batch: 10000/50000 (20%)] Loss: 1.452963
Epoch: 2/10 [Batch: 20000/50000 (40%)] Loss: 1.290976
Epoch: 2/10 [Batch: 30000/50000 (60%)] Loss: 1.135127
Epoch: 2/10 [Batch: 40000/50000 (80%)] Loss: 1.257475
Epoch: 2/10 [Batch: 50000/50000 (100%)] Loss: 1.210437
Epoch: 3/10 [Batch: 10000/50000 (20%)] Loss: 1.235606
Epoch: 3/10 [Batch: 20000/50000 (40%)] Loss: 1.298739
Epoch: 3/10 [Batch: 30000/50000 (60%)] Loss: 1.053390
Epoch: 3/10 [Batch: 40000/50000 (80%)] Loss: 1.167737
Epoch: 3/10 [Batch: 50000/50000 (100%)] Loss: 1.064261
Epoch: 4/10 [Batch: 10000/50000 (20%)] Loss: 0.997144
Epoch: 4/10 [Batch: 20000/50000 (40%)] Loss: 1.166536
Epoch: 4/10 [Batch: 30000/50000 (60%)] Loss: 1.001357
Epoch: 4/10 [Batch: 40000/50000 (80%)] Loss: 0.896586
Epoch: 4/10 [Batch: 50000/50000 (100%)] Loss: 1.036403
Epoch: 5/10 [Batch: 10000/50000 (20%)] Loss: 0.822211
Epoch: 5/10 [Batch: 20000/50000 (40%)] Loss: 1.008645
Epoch: 5/10 [Batch: 30000/50000 (60%)] Loss: 1.085580
Epoch: 5/10 [Batch: 40000/50000 (80%)] Loss: 1.000163
Epoch: 5/10 [Batch: 50000/50000 (100%)] Loss: 0.996731
```

```


Epoch: 6/10 [Batch: 10000/50000 (20%)] Loss: 0.596479
Epoch: 6/10 [Batch: 20000/50000 (40%)] Loss: 0.890222
Epoch: 6/10 [Batch: 30000/50000 (60%)] Loss: 1.112347
Epoch: 6/10 [Batch: 40000/50000 (80%)] Loss: 0.890724
Epoch: 6/10 [Batch: 50000/50000 (100%)] Loss: 1.186765
Epoch: 7/10 [Batch: 10000/50000 (20%)] Loss: 0.753974
Epoch: 7/10 [Batch: 20000/50000 (40%)] Loss: 1.122337
Epoch: 7/10 [Batch: 30000/50000 (60%)] Loss: 0.669767
Epoch: 7/10 [Batch: 40000/50000 (80%)] Loss: 0.670036
Epoch: 7/10 [Batch: 50000/50000 (100%)] Loss: 0.990399
Epoch: 8/10 [Batch: 10000/50000 (20%)] Loss: 0.890396
Epoch: 8/10 [Batch: 20000/50000 (40%)] Loss: 0.645241
Epoch: 8/10 [Batch: 30000/50000 (60%)] Loss: 0.788386
Epoch: 8/10 [Batch: 40000/50000 (80%)] Loss: 0.547218
Epoch: 8/10 [Batch: 50000/50000 (100%)] Loss: 0.750539
Epoch: 9/10 [Batch: 10000/50000 (20%)] Loss: 0.791921
Epoch: 9/10 [Batch: 20000/50000 (40%)] Loss: 0.769445
Epoch: 9/10 [Batch: 30000/50000 (60%)] Loss: 0.877653
Epoch: 9/10 [Batch: 40000/50000 (80%)] Loss: 0.765645
Epoch: 9/10 [Batch: 50000/50000 (100%)] Loss: 0.905367
Epoch: 10/10 [Batch: 10000/50000 (20%)] Loss: 0.690852
Epoch: 10/10 [Batch: 20000/50000 (40%)] Loss: 0.613525
Epoch: 10/10 [Batch: 30000/50000 (60%)] Loss: 0.750566
Epoch: 10/10 [Batch: 40000/50000 (80%)] Loss: 0.722573
Epoch: 10/10 [Batch: 50000/50000 (100%)] Loss: 0.645704

```

```

[ ]: #PLOT THE LEARNING CURVE
      iterations = np.linspace(0,epoch,num_of_batch*epoch)
      plt.plot(iterations, loss_values)
      plt.title('Learning Curve')
      plt.xlabel('Epoch')
      plt.ylabel('Loss')
      plt.grid('on')
      plt.savefig('drive/MyDrive/583_HW7/Q7.png')

```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1

print('Validation accuracy: %.2f%%' %((100*correct)/(total)))
```

Validation accuracy: 68.00%

As proposed in the previous question, with a bigger network (with more parameters), the training loss significantly decreased even with respect to the network in Q6. And it resulted in an 13% increase in the validation accuracy.

## 2.8 Q8

### 2.8.1 Free design

```
[ ]: #DEFINE NEURAL NETWORK MODEL
class CNNQ8(torch.nn.Module):
    def __init__(self):
        super(CNNQ8, self).__init__()
        self.conv1 = torch.nn.Conv2d(3, 32, kernel_size = 3, stride = 1)
        self.conv2 = torch.nn.Conv2d(32, 64, kernel_size = 3, stride = 1)
        self.conv3 = torch.nn.Conv2d(64, 128, kernel_size = 3, stride = 1)
        self.conv4 = torch.nn.Conv2d(128, 128, kernel_size = 3, stride = 1)
        self.conv5 = torch.nn.Conv2d(128, 256, kernel_size = 3, stride = 1)
        self.conv6 = torch.nn.Conv2d(256, 256, kernel_size = 3, stride = 1)
        self.mpool = torch.nn.MaxPool2d(2)
        self.fc1 = torch.nn.Linear(3200, 4096)
        self.fc2 = torch.nn.Linear(4096, 1024)
        self.fc3 = torch.nn.Linear(1024, 10)
        self.relu = torch.nn.ReLU()
        self.sigmoid = torch.nn.Sigmoid()
        self.drop = torch.nn.Dropout(0.1)
        self.conv2bn = torch.nn.BatchNorm2d(64)
        self.conv4bn = torch.nn.BatchNorm2d(128)

    def forward(self, x):
        hidden = self.mpool(self.conv2bn(self.relu(self.conv2(self.relu(self.
→conv1(x))))))
        hidden = self.mpool(self.conv4bn(self.relu(self.conv4(self.relu(self.
→conv3(hidden))))))
        #hidden = self.mpool((self.relu(self.conv6(self.conv5(hidden))))))
        hidden = hidden.view(-1,3200)
        hidden = self.relu(self.fc1(self.drop(hidden)))
        hidden = self.relu(self.fc2(self.drop(hidden)))
        output = self.fc3(hidden)
        return output
```

```
[ ]: #CREATE MODEL
model = CNNQ8()
model.to(device)

# NEW IMPORT
from torch.optim.lr_scheduler import ExponentialLR, ReduceLROnPlateau

#DEFINE LOSS FUNCTION AND OPTIMIZER
learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()
```

```
optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate,
↪weight_decay=1e-5)
scheduler = ExponentialLR(optimizer, gamma=0.1, verbose=True)

# SUMMARY
summary(model, (3,32,32))
```

Adjusting learning rate of group 0 to 1.0000e-03.

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 32, 30, 30]	896
ReLU-2	[-1, 32, 30, 30]	0
Conv2d-3	[-1, 64, 28, 28]	18,496
ReLU-4	[-1, 64, 28, 28]	0
BatchNorm2d-5	[-1, 64, 28, 28]	128
MaxPool2d-6	[-1, 64, 14, 14]	0
Conv2d-7	[-1, 128, 12, 12]	73,856
ReLU-8	[-1, 128, 12, 12]	0
Conv2d-9	[-1, 128, 10, 10]	147,584
ReLU-10	[-1, 128, 10, 10]	0
BatchNorm2d-11	[-1, 128, 10, 10]	256
MaxPool2d-12	[-1, 128, 5, 5]	0
Dropout-13	[-1, 3200]	0
Linear-14	[-1, 4096]	13,111,296
ReLU-15	[-1, 4096]	0
Dropout-16	[-1, 4096]	0
Linear-17	[-1, 1024]	4,195,328
ReLU-18	[-1, 1024]	0
Linear-19	[-1, 10]	10,250
=====		

Total params: 17,558,090

Trainable params: 17,558,090

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 2.42

Params size (MB): 66.98

Estimated Total Size (MB): 69.41

[ ]: #TRAIN THE MODEL

```
model.train()
```

```
epoch = 20
```

```
num_of_batch=np.int(len(train_generator.dataset)/batch_size)
```



```

loss_values = np.zeros(epoch*num_of_batch)
for i in range(epoch):
    if i in [5, 12, 16]:
        scheduler.step()
    for batch_idx, (x_train, y_train) in enumerate(train_generator):
        x_train, y_train = x_train.to(device), y_train.to(device)
        optimizer.zero_grad()
        y_pred = model(x_train)
        loss = loss_fun(y_pred, y_train)
        loss_values[num_of_batch*i+batch_idx] = loss.item()
        loss.backward()
        optimizer.step()
    if (batch_idx+1) % batch_size == 0:
        print('Epoch: {}/{} [Batch: {}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
            i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.
→dataset),
            100. * (batch_idx+1) / len(train_generator), loss.item()))

```

```

Epoch: 1/20 [Batch: 10000/50000 (20%)] Loss: 1.515501
Epoch: 1/20 [Batch: 20000/50000 (40%)] Loss: 1.419433
Epoch: 1/20 [Batch: 30000/50000 (60%)] Loss: 1.297556
Epoch: 1/20 [Batch: 40000/50000 (80%)] Loss: 1.002507
Epoch: 1/20 [Batch: 50000/50000 (100%)] Loss: 1.056283
Epoch: 2/20 [Batch: 10000/50000 (20%)] Loss: 0.911783
Epoch: 2/20 [Batch: 20000/50000 (40%)] Loss: 0.871352
Epoch: 2/20 [Batch: 30000/50000 (60%)] Loss: 0.772256
Epoch: 2/20 [Batch: 40000/50000 (80%)] Loss: 0.931136
Epoch: 2/20 [Batch: 50000/50000 (100%)] Loss: 0.537306
Epoch: 3/20 [Batch: 10000/50000 (20%)] Loss: 0.518227
Epoch: 3/20 [Batch: 20000/50000 (40%)] Loss: 0.774560
Epoch: 3/20 [Batch: 30000/50000 (60%)] Loss: 0.491370
Epoch: 3/20 [Batch: 40000/50000 (80%)] Loss: 0.517129
Epoch: 3/20 [Batch: 50000/50000 (100%)] Loss: 0.425729
Epoch: 4/20 [Batch: 10000/50000 (20%)] Loss: 0.340072
Epoch: 4/20 [Batch: 20000/50000 (40%)] Loss: 0.345048
Epoch: 4/20 [Batch: 30000/50000 (60%)] Loss: 0.406142
Epoch: 4/20 [Batch: 40000/50000 (80%)] Loss: 0.586911
Epoch: 4/20 [Batch: 50000/50000 (100%)] Loss: 0.429337
Epoch: 5/20 [Batch: 10000/50000 (20%)] Loss: 0.192711
Epoch: 5/20 [Batch: 20000/50000 (40%)] Loss: 0.320740
Epoch: 5/20 [Batch: 30000/50000 (60%)] Loss: 0.461353
Epoch: 5/20 [Batch: 40000/50000 (80%)] Loss: 0.394040
Epoch: 5/20 [Batch: 50000/50000 (100%)] Loss: 0.236936
Adjusting learning rate of group 0 to 1.0000e-04.
Epoch: 6/20 [Batch: 10000/50000 (20%)] Loss: 0.133901
Epoch: 6/20 [Batch: 20000/50000 (40%)] Loss: 0.065184
Epoch: 6/20 [Batch: 30000/50000 (60%)] Loss: 0.059485

```

Epoch: 6/20 [Batch: 40000/50000 (80%)] Loss: 0.131645  
 Epoch: 6/20 [Batch: 50000/50000 (100%)] Loss: 0.107893  
 Epoch: 7/20 [Batch: 10000/50000 (20%)] Loss: 0.027405  
 Epoch: 7/20 [Batch: 20000/50000 (40%)] Loss: 0.013605  
 Epoch: 7/20 [Batch: 30000/50000 (60%)] Loss: 0.027041  
 Epoch: 7/20 [Batch: 40000/50000 (80%)] Loss: 0.057117  
 Epoch: 7/20 [Batch: 50000/50000 (100%)] Loss: 0.015393  
 Epoch: 8/20 [Batch: 10000/50000 (20%)] Loss: 0.029922  
 Epoch: 8/20 [Batch: 20000/50000 (40%)] Loss: 0.010794  
 Epoch: 8/20 [Batch: 30000/50000 (60%)] Loss: 0.009067  
 Epoch: 8/20 [Batch: 40000/50000 (80%)] Loss: 0.018730  
 Epoch: 8/20 [Batch: 50000/50000 (100%)] Loss: 0.006877  
 Epoch: 9/20 [Batch: 10000/50000 (20%)] Loss: 0.007529  
 Epoch: 9/20 [Batch: 20000/50000 (40%)] Loss: 0.034778  
 Epoch: 9/20 [Batch: 30000/50000 (60%)] Loss: 0.011480  
 Epoch: 9/20 [Batch: 40000/50000 (80%)] Loss: 0.009583  
 Epoch: 9/20 [Batch: 50000/50000 (100%)] Loss: 0.028255  
 Epoch: 10/20 [Batch: 10000/50000 (20%)] Loss: 0.004036  
 Epoch: 10/20 [Batch: 20000/50000 (40%)] Loss: 0.008029  
 Epoch: 10/20 [Batch: 30000/50000 (60%)] Loss: 0.007810  
 Epoch: 10/20 [Batch: 40000/50000 (80%)] Loss: 0.009471  
 Epoch: 10/20 [Batch: 50000/50000 (100%)] Loss: 0.004101  
 Epoch: 11/20 [Batch: 10000/50000 (20%)] Loss: 0.002351  
 Epoch: 11/20 [Batch: 20000/50000 (40%)] Loss: 0.006347  
 Epoch: 11/20 [Batch: 30000/50000 (60%)] Loss: 0.003042  
 Epoch: 11/20 [Batch: 40000/50000 (80%)] Loss: 0.006736  
 Epoch: 11/20 [Batch: 50000/50000 (100%)] Loss: 0.006071  
 Epoch: 12/20 [Batch: 10000/50000 (20%)] Loss: 0.022980  
 Epoch: 12/20 [Batch: 20000/50000 (40%)] Loss: 0.003312  
 Epoch: 12/20 [Batch: 30000/50000 (60%)] Loss: 0.003654  
 Epoch: 12/20 [Batch: 40000/50000 (80%)] Loss: 0.002607  
 Epoch: 12/20 [Batch: 50000/50000 (100%)] Loss: 0.005876  
 Adjusting learning rate of group 0 to 1.0000e-05.  
 Epoch: 13/20 [Batch: 10000/50000 (20%)] Loss: 0.001097  
 Epoch: 13/20 [Batch: 20000/50000 (40%)] Loss: 0.002119  
 Epoch: 13/20 [Batch: 30000/50000 (60%)] Loss: 0.001296  
 Epoch: 13/20 [Batch: 40000/50000 (80%)] Loss: 0.002892  
 Epoch: 13/20 [Batch: 50000/50000 (100%)] Loss: 0.001395  
 Epoch: 14/20 [Batch: 10000/50000 (20%)] Loss: 0.002301  
 Epoch: 14/20 [Batch: 20000/50000 (40%)] Loss: 0.002234  
 Epoch: 14/20 [Batch: 30000/50000 (60%)] Loss: 0.000866  
 Epoch: 14/20 [Batch: 40000/50000 (80%)] Loss: 0.001604  
 Epoch: 14/20 [Batch: 50000/50000 (100%)] Loss: 0.000824  
 Epoch: 15/20 [Batch: 10000/50000 (20%)] Loss: 0.000675  
 Epoch: 15/20 [Batch: 20000/50000 (40%)] Loss: 0.003404  
 Epoch: 15/20 [Batch: 30000/50000 (60%)] Loss: 0.003100  
 Epoch: 15/20 [Batch: 40000/50000 (80%)] Loss: 0.001102  
 Epoch: 15/20 [Batch: 50000/50000 (100%)] Loss: 0.002950

```


Epoch: 16/20 [Batch: 10000/50000 (20%)] Loss: 0.001383
Epoch: 16/20 [Batch: 20000/50000 (40%)] Loss: 0.001638
Epoch: 16/20 [Batch: 30000/50000 (60%)] Loss: 0.000955
Epoch: 16/20 [Batch: 40000/50000 (80%)] Loss: 0.000966
Epoch: 16/20 [Batch: 50000/50000 (100%)] Loss: 0.002745
Adjusting learning rate of group 0 to 1.0000e-06.
Epoch: 17/20 [Batch: 10000/50000 (20%)] Loss: 0.003343
Epoch: 17/20 [Batch: 20000/50000 (40%)] Loss: 0.003922
Epoch: 17/20 [Batch: 30000/50000 (60%)] Loss: 0.001207
Epoch: 17/20 [Batch: 40000/50000 (80%)] Loss: 0.001179
Epoch: 17/20 [Batch: 50000/50000 (100%)] Loss: 0.001445
Epoch: 18/20 [Batch: 10000/50000 (20%)] Loss: 0.000861
Epoch: 18/20 [Batch: 20000/50000 (40%)] Loss: 0.002336
Epoch: 18/20 [Batch: 30000/50000 (60%)] Loss: 0.001741
Epoch: 18/20 [Batch: 40000/50000 (80%)] Loss: 0.043052
Epoch: 18/20 [Batch: 50000/50000 (100%)] Loss: 0.001661
Epoch: 19/20 [Batch: 10000/50000 (20%)] Loss: 0.001674
Epoch: 19/20 [Batch: 20000/50000 (40%)] Loss: 0.000358
Epoch: 19/20 [Batch: 30000/50000 (60%)] Loss: 0.014443
Epoch: 19/20 [Batch: 40000/50000 (80%)] Loss: 0.000468
Epoch: 19/20 [Batch: 50000/50000 (100%)] Loss: 0.002988
Epoch: 20/20 [Batch: 10000/50000 (20%)] Loss: 0.000919
Epoch: 20/20 [Batch: 20000/50000 (40%)] Loss: 0.000984
Epoch: 20/20 [Batch: 30000/50000 (60%)] Loss: 0.001530
Epoch: 20/20 [Batch: 40000/50000 (80%)] Loss: 0.001313
Epoch: 20/20 [Batch: 50000/50000 (100%)] Loss: 0.000280

```

```

[ ]: #PLOT THE LEARNING CURVE
iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid('on')
plt.savefig('drive/MyDrive/583_HW7/Q8.png')

```



```
[ ]: #TEST THE MODEL
model.eval()
correct=0
total=0

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

    output = model(x_val)
    y_pred = output.argmax(dim=1)

    for i in range(y_pred.shape[0]):
        if y_val[i]==y_pred[i]:
            correct += 1
            total +=1

print('Validation accuracy: %.2f%%' %((100*correct)//(total)))
```

Validation accuracy: 80.00%

With the introduction of Learning Rate Scheduler, Dropout and Batchnorm Layers to the network, it is able to achieve 80% validation accuracy after 20 epochs.